


# TRADITIONAL DATA-DRIVEN SCIENTIFIC DISCOVERY METHODS DO NOT SCALE TO LARGE DATASETS

- 7km GEOS-5 “Nature Run”
- 1 dataset, 3.5 PB
- theoretically: openly accessible
- practically: precomputed pics






# COMMON WORKFLOW FOR SCIENTIFIC DATA ANALYSIS

1. Data Management: Acquisition, Conversion, and Regridding
2. Computation and Analysis
3. Visualization / Comparison

# SPECIFIC TECHNICAL CHALLENGES THAT PREVENT INTERACTIVE SCALING OF DATA- DRIVEN DISCOVERY TO LARGE MODELS

- File formats unsuitable for streaming
- Batch mode data analyses
  - system: submit job and wait
  - algorithm: only final results
- Programming models not progressive
- Server-side analysis does not scale **to large communities**

# WE ADDRESS THE CHALLENGE OF INTERACTIVE EXPLORATION AND ANALYSIS OF MULTI-PETABYTE DATASETS WITHOUT MASSIVE HPC RESOURCES






# INTERACTIVE MULTIRESOLUTION EXPLORATION OF MASSIVE REMOTE DATASETS

(Please show multires\_nature video now)

# METHOD OVERVIEW

- Generic EDSL scripting
- Progressive Runtime environment
- On-demand data reordering




# EMBEDDED DOMAIN SPECIFIC LANGUAGE AND RUNTIME FOR PROGRESSIVE COMPUTATION

- Incremental computation results
  - EDSL supports...
    - abstract data type (location, resolution, format)
    - unordered loops
    - incremental publishing
  - renders current measure of data access irrelevant

# PROGRESSIVE RUNTIME FOR INCREMENTAL SCRIPT EXECUTION

- incremental results
- resolution level
- loop order and parallelization
- server-side processing





# WHY AN EDSL? THE IMPORTANCE OF GENERICITY IN ANALYSIS SCRIPTS

- Genericity facilitates runtime utilization of...
  - incremental execution
  - data format advantages
  - superior loop ordering
  - remote processing
  - parallelization

# EMBEDDED DSL FOR INCREMENTAL COMPUTING

- Minimal extensions to host language
  - ***doPublish, scientific data type, generic loops***
- Example operations of our data type:
  - elementwise combinations (add, subtract, multiply, divide, ...)
  - statistical calculation (average, standard deviation, range, median, ...)
  - domain selection (crop, paste, resize, interleave, ...)

# EDSL SCRIPT FOR INCREMENTAL COMPUTATION OF AVERAGE


```
var output = Array.new();
var i = 0;
unordered(time, [start, end])          // generic loop
{
    var field = Array.read('fieldname', time);
    // critical section (output and i must be updated atomically)
    {{  

        output += (field - output) / (i + 1); // update incremental average  


        i++;
    }}  

    doPublish(); // make intermediate result available
}
```


# EFFECT OF LOOP ORDERING AND STREAMING DATA FORMAT FOR PETA-SCALE CLIMATE DATA ANALYSIS AND VISUALIZATION



# PROGRAMMING MODEL AND RUNTIME SYSTEM ALLOW ALTERNATIVE DATA ORGANIZATION AND PROGRESSIVE COMPUTATIONS



# PROGRAMMING MODEL AND RUNTIME SYSTEM ALLOW ALTERNATIVE DATA ORGANIZATION AND PROGRESSIVE COMPUTATIONS






# COMPARISON OF DIFFERENT LOOP ORDERINGS

(Please show order comparison video now)


# ON-DEMAND DATA REORDERING FOR INTERACTIVE VISUALIZATION AND ANALYSIS OF MASSIVE, DISPARATELY LOCATED DATA

- Interactive access to massive data
  - multiresolution data layouts (IDX)
  - on-demand conversion (operational at LLNL)

# INCREMENTAL MULTIRESOLUTION DATA LOADING




# ON-DEMAND DATA REORDERING FOR OPTIMAL ACCESS OF MASSIVE SPATIOTEMPORAL DATA



1.

Download streaming metadata. Creates and registers a new volume with data server, but does not convert anything.


2.



Client requests data, which is converted on-the-fly to idx multiresolution streaming format and cached for later use.

- Data server first checks idx cache.
- If not cached, requested data is converted can now be loaded from cache.

# EFFECT OF RESOLUTION ON SPEED AND ACCURACY FOR A LARGE COMPUTATION





# DISPARATELY LOCATED DATA CAN BE PROCESSED IN MULTIPLE LOCATIONS AT MOST SUITABLE RESOLUTION

- Eases multi-ensemble analyses
  - automatic regridding (user has control)
  - server-side distributed computation


# RUNTIME: SERVER-SIDE PROCESSING

- Server-side processing:
  - Identical scripting engine as client
  - Remote computation can dramatically reduce data transfer
  - Local vs remote computation specified per script

# USING INTERACTIVE ANALYSIS FOR DATASET ERROR DISCOVERY

correct zonal average

error: every 30 days  
duplicate first day!



# CONCLUSION


- EDSL + Runtime for **interactive, peta-scale** data exploration
- **Incremental results** of dataflow execution
- **On-demand reordering** for optimal data access

# WHERE WE'RE GOING FROM HERE

- Compression of multiresolution data (w/ Peter Lindstrom, LLNL)
  - exploring bit-level precision (e.g., specify 2-bits per item)
- Use abstract computation graph for distribution of computation
- Automatically determine resolution level, loop order, remote, ...
- Web interface
- Docker deployment

# FLEXIBLE COMPRESSION CAPABILITIES: COMBINING DATA REDUCTION IN BOTH RESOLUTION AND PRECISION

- Achieving better compression than each dimension alone



# INTEGRATION WITH ESGF AND EXISTING ANALYSIS TOOLS

- 2013: Using Python SWIG wrappers, UV-CDAT/ViSUS integration
- 2014: Automatic regridding; scripting system is created
- 2015: On-demand conversion of ESGF-hosted data to IDX
- 2016: EDSL formalized; server-side scripting added
- 2017: Integrate CDMS2 module level to enable first class treatment of IDX datasets
- 2017: Multinode incremental server-side EDSL execution
- 2017: Web interface for EDSL-based data analysis and vis



# THANK YOU

- Co-authors: Shusen Liu, Giorgio Scorzelli, Ji-Woo Lee, Peer-Timo Bremer, Valerio Pascucci
- Collaborators: Dean Williams, Sasha Ames, Anthony Hoang, Sam Fries at LLNL
- Funding: DOE NNSA, DOE DREAM, NSF, CCMSC, PIPER, ESGF, LLNL AIMS
- Support and Suggestions: Duong Hoang, Sidharth Kumar, Brian Summa, Amy Gooch, Vidhi Zala

# EMBEDDED DSL AND RUNTIME FOR PROGRESSIVE SPATIOTEMPORAL DATA ANALYSIS AND VIS


- Focus on data analysis, not data management
- Incremental results, interruptible execution, interactive exploration
- Transparent handling of massive, disparately-located data

```
// Computes running average
field = 'TOTSCATAU';           // aerosol scattering
start = query_time;           // current time
width = 720;                  // 720 hours (30 days)


output=Array.New();            // initialize output
var i=0;
unordered(t,[ start , start+width]) // 1d iterator, index t
{
    f=input[field+"?time="+t];      // read field at time t

    // critical section for running average:
    // average and count must be updated atomically
    {{
        output += (f-output)/(i+1); // Welford's method
        i++;
    }}

    doPublish();                  // show current result
}
```



naive method



our method

10% complete

final result